

RUL Prediction as Regression



Say we want to define a RUL-based maintenance policy

How would you tackle that problem?



RUL Prediction as Regression

Let's start from **the simpler formulation** of a RUL-based policy

- We will predict the RUL using a regression approach
- ...And trigger maintenance when the estimated RUL becomes too low, i.e.:

$$f(x, \theta) \leq \varepsilon$$

- f is the regressor, with parameter vector θ
- The threshold ε must account for possible estimation errors

We will focus on the hardest of the four datasets (to reduce training times):

```
In [3]: data_by_src = util.split_by_field(data, field='src')
        dt = data_by_src['train_FD004']
```



We now need to define our training and test data
How do we proceed?



Training and Test Data

We now need to define our training and test data

In a practical setting:

- Some run-to-failure experiments will form the training set
- Others run-to-failure experiments will be used for testing

I.e. we split **whole experiments** rather than individual examples!

Each run-to-failure experiment in our data is associated to a machine

Let's check how many we have:

```
In [4]: print(f'Number of machines: {len(dt.machine.unique())}')
```

```
Number of machines: 249
```

- This is actually a **very large** number
-  In most practical setting, **much fewer** experiments will be available

Training and Test Data

Let's use 75% of the machine for training, the rest for testing

First, we partition the machine indexes:

```
In [5]: tr_ratio = 0.75
        np.random.seed(42)
        machines = dt.machine.unique()
        np.random.shuffle(machines)

        sep = int(tr_ratio * len(machines))
        tr_mcn = machines[:sep]
        ts_mcn = machines[sep:]
```

Then, we partition the dataset itself (via a helper function):

```
In [6]: tr, ts = util.partition_by_machine(dt, tr_mcn)
```



Training and Test Data

Let's have a look at the training data


In [7]:

```
tr
```

Out[7]:

	src	machine	cycle	p1	p2	p3	s1	s2	s3	s4	...	s13	s14	s15	s1
0	train_FD004	461	1	42.0049	0.8400	100.0	445.00	549.68	1343.43	1112.93	...	2387.99	8074.83	9.3335	0.0
1	train_FD004	461	2	20.0020	0.7002	100.0	491.19	606.07	1477.61	1237.50	...	2387.73	8046.13	9.1913	0.0
2	train_FD004	461	3	42.0038	0.8409	100.0	445.00	548.95	1343.12	1117.05	...	2387.97	8066.62	9.4007	0.0
3	train_FD004	461	4	42.0000	0.8400	100.0	445.00	548.70	1341.24	1118.03	...	2388.02	8076.05	9.3369	0.0
4	train_FD004	461	5	25.0063	0.6207	60.0	462.54	536.10	1255.23	1033.59	...	2028.08	7865.80	10.8366	0.0
...
60989	train_FD004	708	180	35.0019	0.8409	100.0	449.44	556.28	1377.65	1148.96	...	2387.77	8048.91	9.4169	0.0
60990	train_FD004	708	181	0.0023	0.0000	100.0	518.67	643.95	1602.98	1429.57	...	2388.27	8122.44	8.5242	0.0
60991	train_FD004	708	182	25.0030	0.6200	60.0	462.54	536.88	1268.01	1067.09	...	2027.98	7865.18	10.9790	0.0
60992	train_FD004	708	183	41.9984	0.8414	100.0	445.00	550.64	1363.76	1145.72	...	2387.48	8069.84	9.4607	0.0
60993	train_FD004	708	184	0.0013	0.0001	100.0	518.67	643.50	1602.12	1430.34	...	2388.33	8120.43	8.4998	0.0

45385 rows × 28 columns



Training and Test Data

...And at the test data


In [8]:

```
ts
```

Out[8]:

	src	machine	cycle	p1	p2	p3	s1	s2	s3	s4	...	s13	s14	s15	s1
321	train_FD004	462	1	41.9998	0.8400	100.0	445.00	548.99	1341.82	1113.16	...	2387.98	8082.37	9.3300	0.0
322	train_FD004	462	2	9.9999	0.2500	100.0	489.05	604.23	1498.00	1299.54	...	2388.07	8125.46	8.6088	0.0
323	train_FD004	462	3	42.0079	0.8403	100.0	445.00	549.11	1351.47	1126.43	...	2387.93	8082.11	9.2965	0.0
324	train_FD004	462	4	42.0077	0.8400	100.0	445.00	548.77	1345.81	1116.64	...	2387.88	8079.41	9.3200	0.0
325	train_FD004	462	5	24.9999	0.6200	60.0	462.54	537.00	1259.55	1043.95	...	2028.13	7867.08	10.8841	0.0
...
61244	train_FD004	709	251	9.9998	0.2500	100.0	489.05	605.33	1516.36	1315.28	...	2388.73	8185.69	8.4541	0.0
61245	train_FD004	709	252	0.0028	0.0015	100.0	518.67	643.42	1598.92	1426.77	...	2388.46	8185.47	8.2221	0.0
61246	train_FD004	709	253	0.0029	0.0000	100.0	518.67	643.68	1607.72	1430.56	...	2388.48	8193.94	8.2525	0.0
61247	train_FD004	709	254	35.0046	0.8400	100.0	449.44	555.77	1381.29	1148.18	...	2388.83	8125.64	9.0515	0.0
61248	train_FD004	709	255	42.0030	0.8400	100.0	445.00	549.85	1369.75	1147.45	...	2388.66	8144.33	9.1207	0.0

15864 rows × 28 columns



Standardization/Normalization

We will use a Neural Network regressor

...Therefore, we need to make the range of each columns more uniform

- We will **standardize** all parameters and sensor inputs:

```
In [9]: trmean = tr[dt_in].mean()
trstd = tr[dt_in].std().replace(to_replace=0, value=1) # handle static fields

ts_s = ts.copy()
ts_s[dt_in] = (ts_s[dt_in] - trmean) / trstd
tr_s = tr.copy()
tr_s[dt_in] = (tr_s[dt_in] - trmean) / trstd
```

- We will **normalize** the RUL values (i.e. our regression target)

```
In [10]: trmaxrul = tr['rul'].max()

ts_s['rul'] = ts['rul'] / trmaxrul
tr_s['rul'] = tr['rul'] / trmaxrul
```



Standardization/Normalization

Let's check the results

In [11]: `tr_s.describe()`

Out [11]:

	machine	cycle	p1	p2	p3	s1	s2	s3
count	45385.000000	45385.000000	4.538500e+04	4.538500e+04	4.538500e+04	4.538500e+04	4.538500e+04	4.538500e+04
mean	582.490955	133.323896	2.894775e-16	1.302570e-16	1.178889e-16	4.664830e-15	2.522791e-15	1.727041e-15
std	71.283034	89.568561	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00
min	461.000000	1.000000	-1.623164e+00	-1.838222e+00	-2.381839e+00	-1.055641e+00	-1.176507e+00	-1.646830e+00
25%	521.000000	62.000000	-9.461510e-01	-1.031405e+00	4.198344e-01	-1.055641e+00	-8.055879e-01	-6.341243e-01
50%	585.000000	123.000000	6.868497e-02	4.154560e-01	4.198344e-01	-3.917563e-01	-6.336530e-01	-4.718540e-01
75%	639.000000	189.000000	1.218855e+00	8.661917e-01	4.198344e-01	6.926385e-01	7.407549e-01	7.495521e-01
max	708.000000	543.000000	1.219524e+00	8.726308e-01	4.198344e-01	1.732749e+00	1.741030e+00	1.837978e+00

8 rows × 27 columns



Regression Model

We can now define a regression model

We will use a feed-forward neural network (MLP):

```
def build_nn_model(input_shape, output_shape, hidden, output_activation='linear'):
    model_in = keras.Input(shape=input_shape, dtype='float32')
    x = model_in
    for h in hidden:
        x = layers.Dense(h, activation='relu')(x)
    model_out = layers.Dense(output_shape, activation=output_activation)(x)
    model = keras.Model(model_in, model_out)
    return model
```

- The `hidden` argument is a list of sizes for the hidden layers
- ...E.g. `hidden = [64, 32]`




Regression Model

We will start with **the simplest possible** Neural Network

...Meaning a **Linear Regressor**!

```
In [15]: nn1 = util.build_nn_model(input_shape=(len(dt_in), ), output_shape=1, hidden=[])  
util.plot_nn_model(nn1)
```

Out [15]:



- We just need to specify that there are no hidden layers
- Why the simplest? As usual, due to Occam's razor



Training

It's useful to define a generic training function

As usual, you can find it in the `util` module:


```
def train_nn_model(model, X, y, loss,
                    verbose=0, patience=10,
                    validation_split=0.0, **fit_params):
    # Compile the model
    model.compile(optimizer='Adam', loss=loss)
    # Build the early stop callback
    cb = []
    if validation_split > 0:
        cb += [callbacks.EarlyStopping(patience=patience,
                                       restore_best_weights=True)]
    # Train the model
    history = model.fit(X, y, callbacks=cb, validation_split=validation_split,
                        verbose=verbose, **fit_params)
    return history
```



Training

We can now train our model

```
In [16]: nn1 = util.build_nn_model(input_shape=(len(dt_in), ), output_shape=1, hidden=[])  
history = util.train_nn_model(nn1, tr_s[dt_in], tr_s['rul'], loss='mse', epochs=30, validation_data=(tr_s[dt_in], tr_s['rul']))  
util.plot_training_history(history, figsize=figsize)
```



Final loss: 0.0142 (training), 0.0108 (validation)




Training

Let's try with a more complex model

```
In [18]: nn2 = util.build_nn_model(input_shape=(len(dt_in), ), output_shape=1, hidden=[32, 32])  
util.plot_nn_model(nn2)
```

Out [18]:




- Now we have two hidden layers
- ...Each with 32 ReLU neurons



Training

Let's check the loss behavior and compare it to Linear Regression

```
In [23]: nn2 = util.build_nn_model(input_shape=(len(dt_in), ), output_shape=1, hidden=[32, 32])  
history = util.train_nn_model(nn2, tr_s[dt_in], tr_s['rul'], loss='mse', epochs=30, validation_data=(tr_s[dt_in], tr_s['rul']))  
util.plot_training_history(history, figsize=figsize)
```



Final loss: 0.0131 (training), 0.0105 (validation)

- There is a modest improvement w.r.t. Linear Regression




Predictions

We can now obtain the predictions and evaluate their quality

```
In [24]: tr_pred = nn2.predict(tr_s[dt_in], verbose=0).ravel() * trmaxrul  
util.plot_pred_scatter(tr_pred, tr['rul'], figsize=figsize)  
print(f'R2 score: {r2_score(tr["rul"], tr_pred)}')
```

R2 score: 0.5380727648735046



What do you think of these results?




Predictions

The results so far are not comforting

...But it's worth seeing what is going on over time:


```
In [28]: stop = 2*1095  
util.plot_rul(tr_pred[:stop], tr['rul'][:stop], figsize=figsize)
```



Predictions

The situation is similar on the test set:

```
In [29]: ts_pred = nn2.predict(ts_s[dt_in], verbose=0).ravel() * trmaxrul  
util.plot_rul(ts_pred[:stop], ts['rul'][:stop], figsize=figsize)
```



Quality Evaluation

Let's try to recap the situation

- Our accuracy is quite poor
- ...Especially for large RUL values

Possible reasons:

- Large RUL value are somewhat scarce on the dataset
- Fault effects become noticeable only after a while



Quality Evaluation

Let's try to recap the situation

- Our accuracy is quite poor
- ...Especially for large RUL values

Possible reasons:

- Large RUL value are somewhat scarce on the dataset
- Fault effects become noticeable only after a while

But perhaps we don't care! Our goal is not a high accuracy

- We just need to stop at the right time
- ...And our model may still be good enough for that

For a proper evaluation, we need a cost model



Cost Model

We will assume that:

We consider one step of operation as our value unit

- ...So we can express the failure cost in terms of operating steps



Cost Model

We will assume that:

We consider one step of operation as our value unit

- ...So we can express the failure cost in terms of operating steps

Every run end with either failure or maintenance:

- Assuming that the failure cost is higher than maintenance cost
- ...We can disregard the maintenance cost



Cost Model

We will assume that:

We consider one step of operation as our value unit

- ...So we can express the failure cost in terms of operating steps

Every run end with either failure or maintenance:

- Assuming that the failure cost is higher than maintenance cost
- ...We can disregard the maintenance cost

A traditional preventive maintenance policy is also available

- We will never trigger maintenance earlier than such policy
- We only gain value if we beat such policy



Cost Model

Slightly more formally:

- One step of operation brings 1 unit of profit
- A failure costs C units more than maintenance
- We only count what happens after s steps

Formally, let x_k be the times series for machine k , and I_k its set of time steps

- The time step when our policy triggers maintenance is given by:

$$\min\{i \in I_k \mid f(x_{ki}) < \varepsilon\}$$

- A failure occurs if:

$$f(x_{ki}) \geq \varepsilon \quad \forall i \in I_k$$



Cost Model

The whole cost formula **for a single machine** will be:

$$\text{cost}(f(x_k), \varepsilon) = \text{op_profit}(f(x_k), \varepsilon) + \text{fail_cost}(f(x_k), \varepsilon)$$

Where:

$$\text{op_profit}(f(x_k), \varepsilon) = -\max(0, \min\{i \in I_k \mid f(x_{ki}) < \varepsilon\} - s)$$

$$\text{fail_cost}(f(x_k), \varepsilon) = \begin{cases} C & \text{if } f(x_{ki}) \geq \varepsilon \quad \forall i \in I_k \\ 0 & \text{otherwise} \end{cases}$$

- s units of machine operation are guaranteed
- ...So we gain over the default policy only if we stop after that
- Profit is modeled as a negative cost

For the total cost, we need to sum over all machines



Cost Model

Normally, we would proceed as follows

- s is determined by the preventive maintenance schedule
- C must be determined by discussing with the customer

In our example, we will derive both from data

First, we collect all failure times

```
In [30]: failtimes = dt.groupby('machine')['cycle'].max()  
failtimes.head()
```

```
Out[30]: machine  
461      321  
462      299  
463      307  
464      274  
465      193  
Name: cycle, dtype: int64
```



Cost Model

Then, we define s and C based on statistics

```
In [31]: print(failtimes.describe())
safe_interval = failtimes.min()
maintenance_cost = failtimes.max()
```

```
count    249.00000
mean     245.97992
std       73.11080
min      128.00000
25%      190.00000
50%      234.00000
75%      290.00000
max      543.00000
Name: cycle, dtype: float64
```

- For the safe interval s , we choose the minimum failure time
- For the maintenance cost C we choose the largest failure time




Threshold Choice

We can then choose the threshold θ as usual

```
In [36]: cmodel = util.RULCostModel(maintenance_cost=maintenance_cost, safe_interval=safe_interval)
th_range = np.arange(-2, 50)
tr_thr = util.opt_threshold_and_plot(tr['machine'].values, tr_pred, th_range, cmodel, figsi:
print(f'Optimal threshold for the training set: {tr_thr}')
```

Optimal threshold for the training set: 5



Evaluation

Let's see how we fare in terms of cost

```
In [37]: tr_c, tr_f, tr_sl = cmodel.cost(tr['machine'].values, tr_pred, tr_thr, return_margin=True)
         ts_c, ts_f, ts_sl = cmodel.cost(ts['machine'].values, ts_pred, tr_thr, return_margin=True)
         print(f'Avg. cost: {tr_c/len(tr_mcn):.2f} (training), {ts_c/len(ts_mcn):.2f} (test)')
```

Avg. cost: -96.51 (training), -106.02 (test)

We can also evaluate the margin for improvement:

```
In [38]: print(f'Avg. fails: {tr_f/len(tr_mcn)} (training), {ts_f/len(ts_mcn)} (test)')
         print(f'Avg. slack: {tr_sl/len(tr_mcn):.2f} (training), {ts_sl/len(ts_mcn):.2f} (test)')
```

Avg. fails: 0.0 (training), 0.0 (test)

Avg. slack: 19.59 (training), 17.79 (test)

- Slack = distance between when we stop and the failure
- The results are actually quite good!

 ... And we also generalize fairly well

Sequence Input in Neural Models



Sequence Input in Neural Models

Feeding more time steps to our NN might improve the results

- Intuitively, sequences provide information about the **trend**
- This may allow a better RUL estimate w.r.t. using only the current state
- E.g. we may gauge how quickly the component is deteriorating

We will try to build a model capable to processing such input

But first, we need to apply a sliding window:

```
In [39]: wlen = 3  
tr_sw, tr_sw_m, tr_sw_r = util.sliding_window_by_machine(tr_s, wlen, dt_in)  
ts_sw, ts_sw_m, ts_sw_r = util.sliding_window_by_machine(ts_s, wlen, dt_in)
```

- This must be done on a **per machine** basis
- Windows **should not mix data** belonging to different machines!



Sliding Window for Multivariate Data

The `sliding_window_by_machine` relies internally on:

```
def sliding_window_2D(data, wlen, stride=1):  
    # Get shifted _tables_  
    m = len(data)  
    lt = [data.iloc[i:m-wlen+i+1:stride, :].values for i in range(wlen)]  
    # Reshape to _add a new axis_  
    s = lt[0].shape  
    for i in range(wlen):  
        lt[i] = lt[i].reshape(s[0], 1, s[1])  
    # Concatenate  
    wdata = np.concatenate(lt, axis=1)  
    return wdata
```

- It's similar to our code for the univariate case
- The output is a tensor with shape `(n_windows, w_len, n_dims)`



Sliding Window for Multivariate Data

Let's look in deeper detail at the returned data structures

`tr_sw` contain the actual sliding window data:

```
In [42]: tr_sw[0]
```

```
Out[42]: array([[ 1.21931469,  0.86619169,  0.41983436, -1.05564063, -0.79621447,
                 -0.70080293, -0.74549387, -1.1386061 , -1.08249848, -0.99389823,
                 -0.11421637, -0.6315044 , -0.67586863, -0.36411574, -0.98910425,
                  0.41889575,  0.08700467,  0.05991388, -0.69502688, -0.63793104,
                 -0.11268403,  0.41983436, -1.03117521, -1.03187757],
                [-0.26962527,  0.41609996,  0.41983436,  0.6926385 ,  0.71397375,
                  0.56288953,  0.29808726,  0.36365649,  0.3710279 ,  0.33249075,
                  0.65388538,  0.56210134, -0.20641916,  0.32893584,  0.33156802,
                  0.41687122, -0.24758681, -0.12925879, -0.69502688,  0.47652818,
                  0.65613725,  0.41983436,  0.35321893,  0.35869109],
                [ 1.21924025,  0.86908928,  0.41983436, -1.05564063, -0.8157647 ,
                 -0.70372248, -0.7109787 , -1.1386061 , -1.08433606, -0.98831315,
                 -0.11380415, -0.64524209, -0.67586863, -0.37335643, -0.99026013,
                  0.41874002, -0.00870947,  0.14931194, -0.69502688, -0.67388133,
                 -0.11268403,  0.41983436, -1.04527086, -1.02276728]])
```



Sliding Window for Multivariate Data

Let's look in deeper detail at the returned data structures

`tr_sw_m` contains the corresponding machine values

```
In [43]: tr_sw_m
```

```
Out[43]: array([461, 461, 461, ..., 708, 708, 708])
```

- The structure is a plain numpy array

`tr_sw_r` contains the RUL values

```
In [44]: tr_sw_r
```

```
Out[44]: array([0.58671587, 0.58487085, 0.58302583, ..., 0.00369004, 0.00184502,  
               0.          ])
```

- Again, the structure is a plain numpy array



1D Convolutions in Keras

The chosen format is ideal for **1D convolutions** in keras

We have a function to build 1D convolutional model in the `util` module

```
def build_cnn_model(input_shape, output_shape, wlen, conv=[], hidden=[], output_activation='linear'):
    model_in = keras.Input(shape=input_shape, dtype='float32')
    x = model_in
    for k in conv:
        x = layers.Conv1D(k, kernel_size=3, activation='relu')(x)
    x = layers.Flatten()(x)
    for k in hidden:
        x = layers.Dense(k, activation='relu')(x)
    x = layers.Dense(output_shape, activation=output_activation)(x)
    model = keras.Model(model_in, x)
    return model
```


- Each convolution in our code will consider 3 time steps
- We need to "flatten" the input before the fully connected layers

1D Convolutions in Keras

Let's build a 1D convolutional model

```
In [45]: nn3 = util.build_cnn_model(input_shape=(wlen, len(dt_in)), output_shape=1, wlen=wlen, conv=
util.plot_nn_model(nn3)
```

Out [45]:




- We have a single convolution with 32 kernels
- Then a hidden layer with 32 ReLU neurons
- ...And finally the output layer



CNN Training

Let's train our CNN

```
In [46]: nn3 = util.build_cnn_model(input_shape=(wlen, len(dt_in)), output_shape=1, wlen=wlen, conv=
history = util.train_nn_model(nn3, tr_sw, tr_sw_r, loss='mse', epochs=25, validation_split=0.1)
util.plot_training_history(history, figsize=figsize)
```



Final loss: 0.0126 (training), 0.0101 (validation)


- We obtained a marginal improvement
- This suggests that considering sequence be modestly useful in this case

Threshold Optimization

Now we can proceed by choosing a threshold

```
In [29]: tr_pred3 = nn3.predict(tr_sw, verbose=0).ravel() * trmaxrul  
ts_pred3 = nn3.predict(ts_sw, verbose=0).ravel() * trmaxrul  
tr_thr3 = util.opt_threshold_and_plot(tr_sw_m, tr_pred3, th_range, cmodel, figsize=figsize)  
print(f'Optimal threshold for the training set: {tr_thr3}')
```

Optimal threshold for the training set: 17



Evaluation

Let's see how the CNN fares in terms of cost

```
In [30]: tr_c3, tr_f3, tr_sl3 = cmodel.cost(tr_sw_m, tr_pred3, tr_thr3, return_margin=True)
         ts_c3, ts_f3, ts_sl3 = cmodel.cost(ts_sw_m, ts_pred3, tr_thr3, return_margin=True)
         print(f'Cost: {tr_c3/len(tr_mcn):.2f} (training), {ts_c3/len(ts_mcn):.2f} (test)')
         print(f'Avg. fails: {tr_f3/len(tr_mcn):.2f} (training), {ts_f3/len(ts_mcn):.2f} (test)')
         print(f'Avg. slack: {tr_sl3/len(tr_mcn):.2f} (training), {ts_sl3/len(ts_mcn):.2f} (test)')
```

```
Cost: -92.10 (training), -105.86 (test)
Avg. fails: 0.01 (training), 0.00 (test)
Avg. slack: 18.98 (training), 16.21 (test)
```

Which is (more or less) on par with our MLP approach, for which we had:

```
In [31]: print(f'Cost: {tr_c/len(tr_mcn):.2f} (training), {ts_c/len(ts_mcn):.2f} (test)')
         print(f'Avg. fails: {tr_f/len(tr_mcn):.2f} (training), {ts_f/len(ts_mcn):.2f} (test)')
         print(f'Avg. slack: {tr_sl/len(tr_mcn):.2f} (training), {ts_sl/len(ts_mcn):.2f} (test)')
```

```
Cost: -98.37 (training), -109.06 (test)
Avg. fails: 0.00 (training), 0.00 (test)
Avg. slack: 17.85 (training), 14.87 (test)
```



Which stresses a simple, but important point...

Time Series and Sequence Input

Just because you are dealing with time series

...Do not assume that sequence input is useful!

- Sequences matter only if the output is correlated with patterns
- ...That involve multiple time steps

In many practical problems

...A single "state" encodes most of the useful information

- You can think of that as sort of Markov property

Therefore, before using sequences, it makes sense to think

Do you expect sequences to provide useful information?

- E.g. is there some kind of inertia?
- ...And does it matter for the considered problem?

