

# Bayesian (Surrogate-Based) Optimization

---



# Bayesian Optimization

We will use an approach known as **Surrogate-Based Bayesian Optimization**

- It is designed to optimize **blackbox functions**
- I.e. functions with an unknown structure, that can only be evaluated

Formally, they address problems in the form:

$$\min_{x \in B} f(x)$$

- Where  $B$  is a box, i.e. a specification of bounds for each component of  $x$
- In our case, the decision variable  $x$  would be  $\theta$
- ...And the function to be optimized would be the cost

The functions are typically assumed to be **expensive to evaluate**



# Why a Surrogate

Since evaluating  $f$  is expensive, it should be done **infrequently**

The main trick to achieve this is using a **surrogate model**

- After each evaluation we train a **Machine Learning model**
- ...Then we perform optimization **on the ML model**
- ...Since it can be evaluate much more quickly

The process is usually start by sampling a few random points

## This is where the name stems from


- Since we use the ML model instead of the function, we call it a **surrogate**
- Moreover, we optimize over **prior** information (i.e. the current model)
- ...And we refine the model based on the evaluation (**posterior**)
- Hence we call it **Bayesian Optimization**



# A Running Example

Let's assume we want to minimize the following function over  $[0, 1]$

```
In [2]: bbf = lambda x: (0.5 - x**2) * np.sin(2 * 4 * np.pi * x)
xrange = np.linspace(0, 1, 1000)
target = pd.Series(index=xrange, data=bbf(xrange))
util.plot_series(target, figsize=figsize)
```




- There multiple local minima, and the global minimum is at  $\simeq 0.19$



# A Running Example

Let's start by sampling a few points at random

```
In [3]: np.random.seed(42)
xtr = np.sort(np.random.random(4))
ytr = bbf(xtr)
util.plot_series(target, figsize=figsize)
plt.scatter(xtr, ytr, color='tab:orange');
```



- Using **only the orange points** we need to train a model
- ...That gives us a good idea for a new point to evaluate

**Which properties should our surrogate model have?**



# Properties of a Good Surrogate

Our surrogate model should



# Properties of a Good Surrogate

## Our surrogate model should

Approximate very accurately all evaluated points

- Assuming the function is deterministic
- ...The available evaluations are exact values





# Properties of a Good Surrogate

## Our surrogate model should

Approximate very accurately all evaluated points

- Assuming the function is deterministic
- ...The available evaluations are exact values

Reflect our confidence level on unexplored regions

- If we have few samples in a certain region
- ...We might want to search there just to see what the function looks like

**Can you think of a ML model with these properties?**



# Gaussian Process Surrogates

## Gaussian Processes check all the boxes!

- They can interpolate very well known measurements
- They provide a confidence level that decays with distance from observations

## Let's try to train a simple GP for our example

```
In [4]: kernel = RBF(0.01, (1e-3, 1e3)) + WhiteKernel(1e-3, (1e-6, 1e-2))
gpr = GaussianProcessRegressor(kernel=kernel, n_restarts_optimizer=9, normalize_y=True)
gpr.fit(xtr.reshape(-1, 1), ytr);
gpr.kernel_
```

```
Out[4]: RBF(length_scale=0.0792) + WhiteKernel(noise_level=0.000126)
```


- We use an **RBF kernel** to capture the distance-based correlation
- We also use a **white noise kernel** to avoid numerical instability
- ...But we keep it at a low value since the target function is deterministic



# Gaussian Process Surrogate

## Let's inspect our Gaussian Process Surrogate

```
In [5]: pmean, pstd = gpr.predict(xrange.reshape(-1, 1), return_std=True)
util.plot_gp(target=target, figsize=figsize, pred=pd.Series(index=xrange, data=pmean),
             std=pd.Series(index=xrange, data=pstd), samples=pd.Series(index=xtr, data=ytr))
```



- All known points are interpolated (almost) exactly
- ...And the confidence intervals behave in an intuitive fashion




# What to Optimize?

Now we need to search over the surrogate model

This is the same as choosing **which function to optimize**

```
In [6]: pmean, pstd = gpr.predict(xrange.reshape(-1, 1), return_std=True)
util.plot_gp(target=target, figsize=figsize, pred=pd.Series(index=xrange, data=pmean),
             std=pd.Series(index=xrange, data=pstd), samples=pd.Series(index=xtr, data=ytr))
```



Which area does it make sense to explore, and why?

# Acquisition Function

We need to account for both the **predictions** and their **confidence**

- Area with **low predictions** are promising
- ...But so are also areas with **high confidence**

**This issue is solved in SBO by optimizig an acquisition function**

...Which should balance **exploration** and **exploitation**.

- Examples include the Probability of improvement, the Expected Improvement
- ...And the Lower/Upper confidence bound

**We will use the Lower Confidence Bound, which is given by:**

$$LCB(x) = \mu(x) - Z_{\alpha}\sigma(x)$$


- Where  $\mu(x)$  is the predicted mean,  $\sigma(x)$  is the predicted standard deviation
- ...And  $Z_{\alpha}$  is multiplier for a  $\alpha\%$  Normal confidence interval



# Lower Confidence Bound

Let's see an example in our case with  $Z_\alpha = 2.5$

```
In [7]: lcb = pmean - 2.5 * pstd  
util.plot_gp(target=target, figsize=figsize, pred=pd.Series(index=xrange, data=pmean), std=pstd,  
plt.plot(xrange, lcb, color='tab:red');
```




- We can then optimize via any method applicable to our surrogate
- E.g. Nelder-Mead, Mathematical Programming, or even simple grid search



# Lower Confidence Bound

Let's see which point we would choose in our case

```
In [8]: best_idx = np.argmin(lcb)
util.plot_gp(target=target, figsize=figsize, pred=pd.Series(index=xrange, data=pmean), std=
plt.plot(xrange, lcb, color='tab:red');
plt.scatter(xrange[best_idx], lcb[best_idx], marker='x', color='tab:red');
```



- The  $x$  value with the best acquisition function is highlighted with a red "x"



# Updating the Surrogate

## Now we update our surrogate model

First, we evaluate  $f$  for the new point and grow our training set:

```
In [9]: xtr2 = np.hstack((xtr, [xrange[best_idx]]))  
        ytr2 = np.hstack((ytr, [bbf(xrange[best_idx])]))
```

Then we can retrain our Gaussian Process:

```
In [10]: gpr2 = GaussianProcessRegressor(kernel=kernel, n_restarts_optimizer=9, normalize_y=True)  
         gpr2.fit(xtr2.reshape(-1, 1), ytr2);  
         gpr2.kernel_
```

```
Out[10]: RBF(length_scale=0.0999) + WhiteKernel(noise_level=2.46e-06)
```

- Then we should optimize the acquisition function again
- ...But we will limit ourselves to showing the updated predictions






# Updating the Surrogate

Here are the estimates for the update surrogate

...Together with the acquisition function and the next iterate

```
In [11]: pmean2, pstd2 = gpr2.predict(xrange.reshape(-1, 1), return_std=True)
lcb2 = pmean2 - 2.5 * pstd2
best_idx2 = np.argmin(lcb2)
util.plot_gp(target=target, figsize=figsize, pred=pd.Series(index=xrange, data=pmean2), std=pstd2,
plt.plot(xrange, lcb2, color='tab:red');
plt.scatter(xrange[best_idx2], lcb2[best_idx2], marker='x', color='tab:red');
```



# Surrogate-Based Bayesian Optimization

## Let's review the general method

- Given a collection  $\{\mathbf{x}_i, y_i\}_i$  of evaluated points
- ...We train a surrogate-model  $\hat{f}$  for  $f$

## Then we proceed as follows:

- We optimize an acquisition function  $a_{\hat{f}}(\mathbf{x})$  to find a value  $\mathbf{x}'$
- We evaluate  $y' = f(\mathbf{x}')$
- If  $y'$  is better than the current optimum  $f(\mathbf{x}^*)$ :
  - Then we replace  $\mathbf{x}^*$  with  $\mathbf{x}'$
- We expand our collection of measurements to include  $(\mathbf{x}', y')$
- We retrain  $\hat{f}$
- We repeat until a termination condition is reached



# A Few Considerations

## Different Bayesian optimization algorithms:

- Make use of different surrogate models
- Rely on different criteria for choosing  $\mathbf{x}'$
- Strike different trade-offs in terms of number of (expensive) evaluations of  $f$
- ...And the quality of the obtained solutions

**For more information, see (e.g.) [this tutorial](#)**

## In practice, you don't have to code from scratch

...Since multiple libraries are available, like:

- The [scikit-optimize package](#) (crude, reasonably fast, unfortunately unmaintained)

- The [bayesian-optimization python module](#) (more stable, but also slower)



- The [RBFOpt solver](#) (stable, fast, more complex in terms of requirements)

# SBO for Threshold Calibration

---



# Back to Our Motivating Problem

We will use SBO to tackle our policy definition problem

$$\begin{aligned} \operatorname{argmin}_{\varepsilon} \sum_{k \in K} \operatorname{cost}(f(x_k, \theta^*), 1/2) \\ \text{s.t.: } \theta^* = \operatorname{argmin}_{\theta} L(f(x_k, \lambda), \mathbb{1}_{y_k \geq \varepsilon}) \end{aligned}$$

Here's our plan:

- We need to optimize over  $\varepsilon$
- Our goal is minimizing the cost
- Computing the cost requires to re-define the classes
- ...And therefore to repeat training

**Our implementation will be based on bayesian-optimization**



# The Black Box Function

**As a first step, we need to define our black box function**

We will use a function class (in the `util` module) with this structure:

```
class ClassifierCost:
    def __init__(self, machines, X, y, cost_model, init_epochs=20, inc_epochs=3):
        ...

    def __call__(self, params):
        ...
```

- In the constructor, we provide parameters that are fixed during optimization
- In the `__call__` method, we retrain the model and evaluate the cost
- The `__call__` method is executed when we try to invoke an object of this class
- ...Meaning that we can treat an object of this class as a normal function



# The Black Box Function

It is worth having a deeper look at the `__call__` method

```
def __call__(self, params):  
    theta = params[0] # There is only one parameter to optimize  
    lbl = (self.y >= theta) # Redefine classes  
    # Determine the number of epochs and retrain  
    epochs = self.init_epochs if not self.is_init else self.inc_epochs  
    self.is_init = True  
    train_nn_model(self.nn, self.X, lbl, loss='binary_crossentropy', epochs=epochs,  
                   verbose=0, patience=10, batch_size=32, validation_split=0.2)  
    ...
```

- At each execution we redefine the classes
- We use warm starting to make the process faster
- Each training attempt after the first uses only a few epochs



# The Black Box Function

It is worth having a deeper look at the `__call__` method

```
def __call__(self, params):  
    ...  
    self.stored_weights[theta] = self.nn.get_weights() # Store weights  
    # Evaluate cost  
    pred = np.round(self.nn.predict(self.X, verbose=0).ravel())  
    cost, fails, slack = self.cost_model.cost(self.machines, pred, 0.5, return_margin=True)  
    return cost
```

- We store the weights in a dictionary for later retrieval
- We need this to rebuild the optimal network once optimization is over
- Finally, we evaluate the cost
- The actual code in `util` also prints some information





# The Black Box Function

We can build an object in the usual way

```
In [13]: ccf = util.ClassifierCost(machines=tr['machine'], X=tr_s[dt_in], y=tr['rul'], cost_model=cme
```

...But since it is a function, we can **invoke** it:

```
In [14]: ccf(20, verbose=True)
```

```
eps: 20.00, avg. cost: -88.96, avg. fails: 0.00, avg. slack: 27.48
```

```
Out[14]: 16547
```

- The function returns the negated cost
- ...Since **bayesian-optimization** is designed for maximization



# Running the Solver

Now we can define our box constraints and run the optimization process

```
In [15]: pbounds = {'eps': (1, 20)} # Box constraints
optimizer = BayesianOptimization(f=ccf, pbounds=pbounds, random_state=42)
optimizer.maximize(init_points=3, n_iter=10)
```

iter	target	eps
1	1.871e+04	8.116
2	1.657e+04	19.06
3	1.73e+04	14.91
4	-5.761e+0	1.192
5	1.815e+04	11.33
6	1.632e+04	17.07
7	1.812e+04	9.54
8	1.764e+04	13.09
9	1.641e+04	20.0
10	1.827e+04	7.082
11	1.814e+04	7.599
12	1.909e+04	5.732
13	1.586e+04	6.213

- The implementation is very close to what we have showed
- The results will be a bit noisy, since training is stochastic

# Retrieve the Results

We can access the best  $\epsilon$  value from a result data structure

```
In [16]: print(optimizer.max)
best_eps = optimizer.max['params']['eps']

{'target': 19091.0, 'params': {'eps': 5.731594133424523}}
```

We will use it to retrieve the weights of the best network:

```
In [17]: nn = keras.models.clone_model(ccf.nn)
nn.set_weights(ccf.stored_weights[best_eps])
```

- This is possible since we stored all the tested weights in our class



# Evaluate the Classifier

Finally, we can evaluate our classifier

```
In [18]: tr_pred = np.round(nn.predict(tr_s[dt_in], verbose=0).ravel())
         ts_pred = np.round(nn.predict(ts_s[dt_in], verbose=0).ravel())

         tr_c, tr_f, tr_sl = cmodel.cost(tr['machine'].values, tr_pred, 0.5, return_margin=True)
         ts_c, ts_f, ts_sl = cmodel.cost(ts['machine'].values, ts_pred, 0.5, return_margin=True)
         print(f'Avg. cost: {tr_c/len(tr_mcn)} (training), {ts_c/len(ts_mcn)} (test)')
         print(f'Avg. fails: {tr_f/len(tr_mcn)} (training), {ts_f/len(ts_mcn)} (test)')
         print(f'Avg. slack: {tr_sl/len(tr_mcn):.2f} (training), {ts_sl/len(ts_mcn):.2f} (test)')
```

```
Avg. cost: -102.63978494623656 (training), -113 (test)
Avg. fails: 0.0 (training), 0.0 (test)
Avg. slack: 13.42 (training), 11.06 (test)
```

- The performance is now better than the regression approach!



# AutoML

## Many ML models have hyper parameters!

...And tuning them may sometimes improve the performance

- The problem is that tuning multiple parameters may be complicated
- ...And every training attempts is expensive

## This makes hyper-parameter tuning a perfect application for SBO

...And other similar approaches. A few libraries you might have heard of:

- Hyperopt
- Optuna

## In recent years the concept has been generalized to AutoML

...Where we can start chanking the architecture and model type, too!

- It's a big topic (and big techs have some available SW solutions)
- A good starting reference is this web site

